


ADAM: A method for stochastic optimization

Kingma & Ba, 2014 (158k citations)

Mattea Busato, Stefano Mauloni, Daniyar Zakarin

16th November, 2023

Brief Recap on Neural Networks:



The **loss function** is a function of the input data, the network weights and the expected output.

The loss function therefore **defines the task** for which the model is intended and the weights **need to be trained** to fit the task.

Gradient descent variants:

$$\theta_t = \theta_{t-1} - \eta * \frac{\partial L}{\partial \theta_{t-1}}$$

(N total datapoints)



N data points



Classic Gradient Descent



1 data point



Stochastic Gradient Descent



k data points



Mini Batch SGD

Depending on the amount of data, we make a **trade-off** between the **accuracy** of the weights' update and the **computational time** it takes to perform an update.


SGD with Momentum:

Update rule:
$$\begin{cases} \theta_t = \theta_{t-1} + v_t \\ v_t = \rho v_{t-1} - \eta g(\theta_{t-1}) \end{cases} \Rightarrow \theta_t = \theta_{t-1} + \rho v_{t-1} - \eta g(\theta_{t-1})$$

We introduce **momentum** ρ :

1. It helps accelerate SDG in the relevant direction
2. It smooths out the trajectory by mitigating the stochasticity of the motion of θ in parameter space

The choice of learning rate η is made a bit **less crucial** by the fact that the training procedure is **adaptive** to the particular loss landscape.




AdaGrad (Adaptive Gradient):

Update rule:
$$\theta_{i,t} = \theta_{i,t-1} - \frac{\eta}{\varepsilon + \sqrt{\sum_{j=1}^{t-1} g(\theta_{i,j})^2}} * g(\theta_{i,t-1})$$

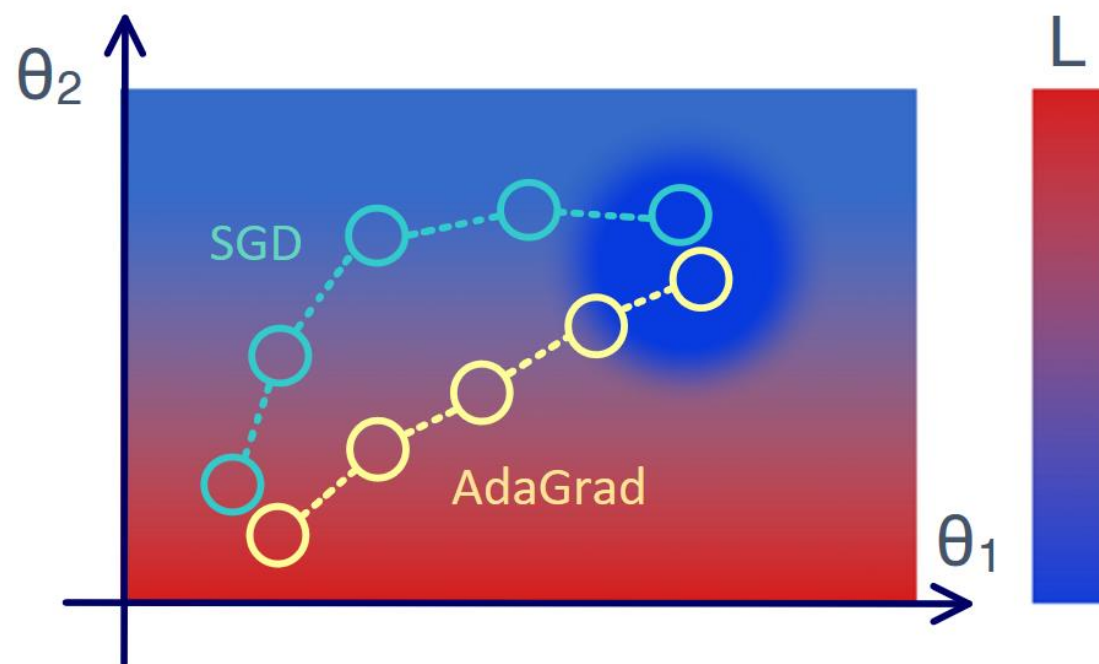
To simplify the notation:
$$\begin{cases} \theta_{i,t} = \theta_{i,t-1} - \frac{\eta}{\varepsilon + \sqrt{v_{i,t}}} * g(\theta_{i,t-1}) \\ v_{i,t} = v_{i,t-1} + g(\theta_{i,t-1})^2 \end{cases}$$

AdaGrad adapts the learning rate of **each parameter** differently as training progresses.



AdaGrad (Adaptive Gradient):

The **idea** is that if a parameter has changed significantly, it must have made a lot of progress towards the target and if it has not changed much, it should keep getting updated with greater emphasis.



RMSProp (Root Mean Square Propagation)

Update rule:

$$\begin{cases} \theta_t = \theta_{t-1} - \frac{\eta}{\varepsilon + \sqrt{v_t}} * g(\theta_{t-1}) \\ v_t = \beta v_{t-1} + (1 - \beta) g(\theta_{t-1})^2 \end{cases}$$

We introduce a **discount parameter** $\beta \in [0, 1]$:

1. It controls how much of the previous term v_{t-1} is remembered
2. It allows the scaling down and scaling up of the learning rate

We retain the benefits of the **decaying learning** rate without the risk of suffering a permanently decayed rate

What is ADAM?

- ADAM is a stochastic gradient-descent optimization method.
- It consists, as we have seen with other methods, in updating weights based on the value of the gradient computed at each time step.

Characteristics:

1. Straightforward to implement
2. Computationally efficient
3. Uses little memory
4. Adapts to the parameters landscape

What is ADAM?

- ADAM is a stochastic gradient-descent optimization method.
- It consists, as we have seen with other methods, in updating weights based on the value of the gradient computed at each time step.

Requirements:

- Stochastic objective function
- Function must be differentiable wrt its parameters θ

Notation

- θ = weights
- v = velocity
- ρ, η = fixed parameters
- $g(\theta) = \nabla_{\theta} f(\theta)$ = gradient of f wrt θ

Update rule of “Classic” SGD with momentum:

$$\begin{cases} \theta_t = \theta_{t-1} + v_t \\ v_t = \rho v_{t-1} - \eta g(\theta_{t-1}) \end{cases} \Rightarrow \theta_t = \theta_{t-1} + \rho v_{t-1} - \eta g(\theta_{t-1})$$

Additional Notation

- β_1, β_2 = discount parameters (how much of the previous is remembered)
- m = first moment estimate
- v = second moment estimate
- ϵ = numerical stabilizer (avoids divisions by zero)

Update rule of ADAM:

- 1st moment vector:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g(\theta_{t-1})$$

- 2nd moment vector:

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g(\theta_{t-1})^2$$

$$\Rightarrow \theta_t = \theta_{t-1} - \frac{\alpha}{\epsilon + \sqrt{\hat{v}_t}} \hat{m}_t$$

Why \hat{v}_t and \hat{m}_t ?

The algorithm computes **exponential moving** (not just decaying) averages of the gradient (m_t) and the squared gradient (v_t).

We need to initialize θ_0 , usually as a vector of zeros. Hence, m_t and v_t will be **biased towards $\mathbf{0} \in \mathbb{R}^n$** (initialization bias).

m_t and v_t are biased moment estimates. Bias is stronger:

- At initial timesteps (small t)
- When decay rates are small (β_1, β_2 close to 1)

To solve this, ADAM introduces **bias-corrected** 1st and 2nd moment estimates:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

NB: They are magnification of the standard estimates for small t , while they tend to be equal to the biased as t becomes larger and larger

Algorithm

Good default settings: $\alpha = 0.1$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$

Require: α , β_1 , β_2 , $f(\theta)$, θ_0

Require: $m_0 \leftarrow 0$, $v_0 \leftarrow 0$, $t \leftarrow 0$

While θ_t not converged do:

$$t \leftarrow t + 1$$

$$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_t + (1 - \beta_2) g_t^2$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$$\theta_t = \theta_{t-1} - \frac{\alpha \hat{m}_t}{\epsilon + \sqrt{\hat{v}_t}}$$

Return θ_t

Update Rule: 3 Properties

1: Boundedness

Assuming $\epsilon = 0$, the effective stepsize in parameter space is:

$$|\Delta_t| = \alpha \frac{\widehat{m}_t}{\sqrt{\widehat{v}_t}}$$

It is **bounded** by two upper bounds:

$$\begin{cases} |\Delta_t| \leq \alpha \frac{(1 - \beta_1)}{\sqrt{(1 - \beta_2)}} & \text{if } (1 - \beta_1) > \sqrt{(1 - \beta_2)} \\ |\Delta_t| \leq \alpha & \text{if } (1 - \beta_1) \leq \sqrt{(1 - \beta_2)} \end{cases}$$

Since the first happens only in very peculiar cases (sparse gradients), in general we can assume $|\Delta_t| \lesssim \alpha$

NB: It establishes a trust region outside of which the current gradient does not provide sufficient information. Hence, **we never move outside of it:**

2: Invariance to rescaling of the gradients

Multiplying g by a factor c will make no difference in the magnitude of the effective stepsize.

Rescaling the gradients g by a factor c :

$$\mathbf{m}'_t = \beta_1 \mathbf{m}'_{t-1} + (1 - \beta_1) c g_{t-1} = \beta_1 c \mathbf{m}_{t-1} + (1 - \beta_1) c g_{t-1} = c \mathbf{m}_t$$

$$\mathbf{v}'_t = \beta_2 \mathbf{v}'_{t-1} + (1 - \beta_2) (c g_{t-1})^2 = \beta_2 c^2 \mathbf{v}_{t-1} + (1 - \beta_2) c^2 (g_{t-1})^2 = c^2 \mathbf{v}_t$$

$$|\Delta_t|' = \alpha \frac{\widehat{m}'_t}{\sqrt{\widehat{v}'_t}} = \alpha \frac{m'_t}{\sqrt{v'_t}} \frac{\sqrt{1 - \beta_2^t}}{1 - \beta_1^t} = \alpha \frac{c m_t}{\sqrt{c^2 v_t}} \frac{\sqrt{1 - \beta_2^t}}{1 - \beta_1^t} = \alpha \frac{m_t}{\sqrt{v_t}} \frac{\sqrt{1 - \beta_2^t}}{1 - \beta_1^t} = |\Delta_t|$$

3: “Cautiousness”

We define signal-to-noise ratio (SNR) $:= \frac{\widehat{m}_t}{\sqrt{\widehat{v}_t}}$

$$\text{SNR} \approx 0 \Rightarrow \Delta_t \approx 0$$

It is desirable, since this means that there is uncertainty whether the true direction of the gradient corresponds to the one of \widehat{m}_t .

This usually happens near a point of minimum, so this prevents jumping around it.

Initialization Bias Correction

Derivation of v_t (analogous for m_t)

1: rewrite v_t as $v_t = (1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} \cdot g_i^2$

2: algebra:

$$\begin{aligned}\mathbb{E}(v_t) &= \mathbb{E}\left((1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} \cdot g_i^2\right) = \mathbb{E}(g_t^2)(1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} + \zeta \\ &= \mathbb{E}(g_t^2)(1 - \beta_2^t) + \zeta \quad \text{bias}\end{aligned}$$

$\zeta = 0$ if the gradient is stationary (namely, constant), otherwise it is still very close to zero since gradients far in the past have small values. **We ignore it.**

Initialization Bias Correction

Example: (with \mathbf{m}_t)

$t = 1$:

$$m_1 = \beta_1 m_0 + (1 - \beta_1) g_1$$

We take out $\beta_1 m_0$ and divide by $(1 - \beta_1)$ to get the same expectation as g_1 ,
This yields:

$$\hat{m}_1 = \frac{m_1 - \beta_1 m_0}{(1 - \beta_1)}, \quad \mathbb{E}(\hat{m}_1) = \mathbb{E}(g_1)$$

But $\beta_1 m_0 = 0$, thus:

$$\hat{m}_1 = \frac{m_1}{(1 - \beta_1)}$$

Note that in this trivial case $\zeta = 0$.

Does it actually converge?


Yes (proved in the paper), assuming bounded gradients.

Not always in practice. It depends at a great extent to the choice of β_2 .


In general, we take β_2 large enough (very close to 1) to ensure convergence.

Power of Momentum: Acceleration


Gradient Descent




Gradient Descent with Momentum




ADAM




Power of Momentum: Noise Cancelling




Second moment helps with sparsity



Second moment helps with sparsity




Logistic Regression: MNIST




- Authors use L2-Regularized Multi Class Logistic regression on MNIST image dataset
- $1 / \sqrt{t}$ learning rate decay was used in agreement with theoretical convergence result

Logistic Regression: IMDB reviews

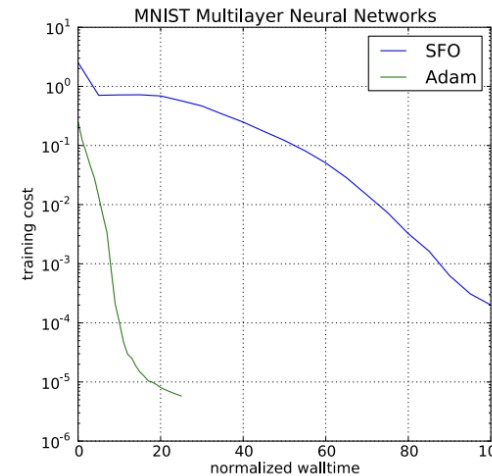



- IMDB movie reviews were preprocessed to bag-of-words (BoW) feature vectors including 10000 most frequent words
- Adagrad performs well compared to SGD on sparse data and sparse gradients
- Notice blue & pink lines are smoother

Feed-forward Neural Networks





(a)



(b)

- The Sum of Functions (SFO) method quasi-Newton method that works with minibatches of data (Sohl-Dickstein et al., 2014)

Convolutional Neural Networks




- Second momentum vanishes
- Adam does not require handpicked layer learning rates

ADAM after 2014

- Attention is All you Need, 2017
- U-Net: Convolutional Networks for Biomedical Image Segmentation, 2015
- BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding, 2019
- An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale, 2020

"ADAM does well where others fail, although not without drawbacks..." © Alex from Stack Exchange

People behind ADAM



Thanks

Mattea Busato, Stefano Mauloni, Daniyar Zakarin